

## Universidad de León

### Bachelor Degree on Computer Science and Engineering

#### Course on Computer Networks

## Weekly Homework no. 5 (WH5-Practice)

All rights reserved © 2013-2020 by José María Foces Morán and José María Foces Vivancos

**Published on:** 3<sup>rd</sup> - May- 2020

**Submission date:** 14<sup>th</sup>-May-2020

**Submit via:** [foces.informatica.unileon@gmail.com](mailto:foces.informatica.unileon@gmail.com)

**Format:** Only pdf format is accepted. Include your name and ID in the document.

### --- Study Guide ---

1. **To be read carefully:** Finally, these weekly homework (WH) assignments do count to the overall final grade according to the official Course Guide updated on 22<sup>nd</sup>-April. The Weekly Homework is a valuable means for you to remain fully involved with the course while away from campus.
2. At all times, have the textbook by Peterson & Davie at hand. Most of the material that we have taught so far belongs in book chapters 1 and 2. Find the 6<sup>th</sup> edition to the book, here:

<https://github.com/SystemsApproach/book/releases/download/v6.1/book.pdf>

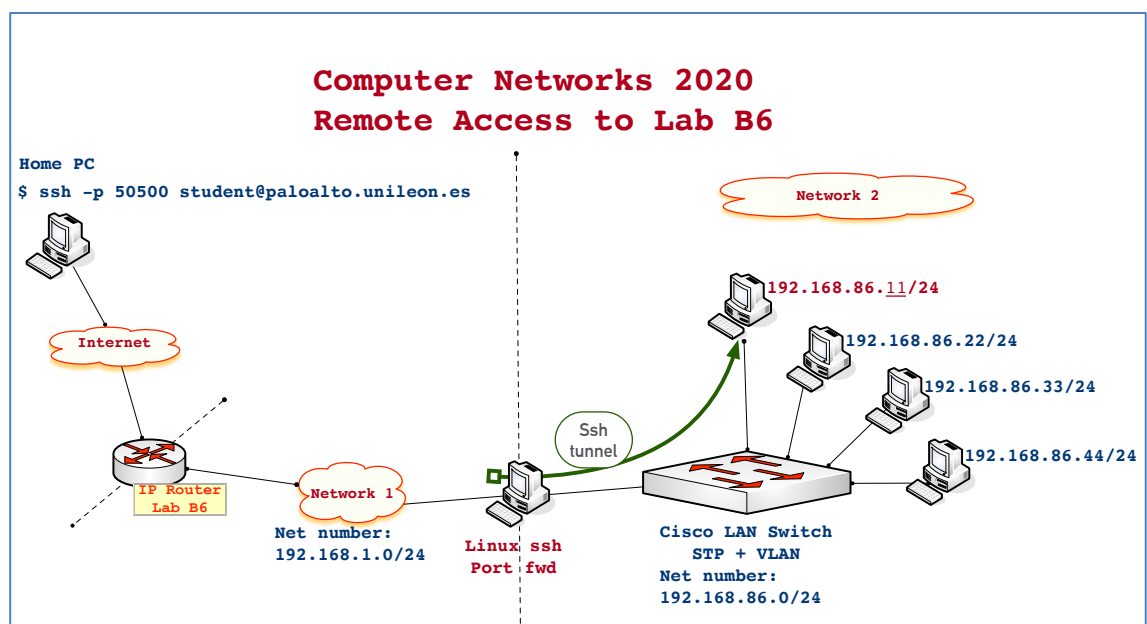
3. A valuable resource as you undertake the practice exercises contained in **WH5-Practice** is the practices that we did the past academic year. Source code and guiding explanations can be found in the practice scripts.

<http://paloalto.unileon.es/cn/>

4. If you need assistance, contact me via the email given above; I recommend that you send your comments and your questions to the **class forum** in the Agora. While RemLabB6 is open, usually I'll be listening to your questions on the CN Chat, which results very practical for obtaining my assistance with the practices.
5. **Make sure that you fully understand the procedures to access the remote network** in Lab B6 which are explained in WH3, in section "[Instructions for accessing Lab B6 Remote Network](#)". When you access RemLabB6 by following the provided instructions and using your login

and password, you end up logging in host **192.168.86.11** in Network 2 (See fig. 1). The practices included in this WH entail your accessing the other hosts in Network 2 with ssh. The user name to be used in these cases is **administrator** and the password is Lab B6's general password: "**19xxdpq16**". Note that you have root privileges in the latter hosts (**192.168.86.22**, **192.168.86.33**, **192.168.86.44**), consequently you must exercise due care when you use them, so that we keep them functional for the whole class.

6. Host 192.168.86.11 in Network 2 is your *trampoline* into Network 2. Every time you access Network 2 via `ssh -p 50500 paloalto.unileon.es`, the internal host that receives your login is 192.168.86.11; this latter host is your *trampoline* to the remaining hosts in Network 2 (**192.168.86.22**, **192.168.86.33**, **192.168.86.44** and others that we will announce in due time).
7. In the present WH, you will write a program in C that accesses the Ethernet directly. Programs that access the network or datalink layers directly require a Linux capability known as *Raw Socket Capability*. In this practice, again, you will scale the capabilities of your programs to be authorized to access Ethernet's service interface consists of sending your program's full path name to a fifo set by user administrator.
8. These these practices offer you training you need for the on-line final exam (Theory and Practices), consequently do these practices attentively, thinking each exercise out and resolve your doubts with your study and by contacting me if necessary.



**Figure1.** Host 192.168.86.11 is your *trampoline* into Network 2. From that host you can open as many ssh sessions with the other hosts (22, 33 and 44) as you need

9. Contact me if you need that I open RemLabB6 at times other than those announced (Tuesdays and Thursdays) or if you need assistance:

[foces.informatica.unileon@gmail.com](mailto:foces.informatica.unileon@gmail.com)

or

the CN Chat in [agora.unileon.es](https://agora.unileon.es)

## Weekly Homework (WH5-Practice)

*This section is to be included in your homework submission. These exercises will be assessed. You must submit your original work and cite sources in case you used some.*

### Base documentation

- **Class notes, presentations and questionnaires about the Internet Architecture** (Lectures from February and early March).
- **Refresher about Internet Architecture (IA) and Interfaces.** In the lectures referenced just above, we introduced the Internet Architecture (IA). The essential element of IA is a hierarchy of four protocol layers and interfaces to layers 1, 2 and 3. Each interface permits a programmer to use the services provided by the accessed layer's protocol. You, the programmer simply invoke some `send()`<sup>1</sup> or `receive()` function call from the sockets library. The rest of the communication processes are completely transparent to you, the programmer; those *processes* are the responsibility of the socket interface alongside with the underlying protocol stack.

Each of the lower three layers in the IA offers a socket interface, consequently you can make network programs that choose a socket interface to access one of the following layers:


- **Interface to Layer 3.** These sockets allow programming against the UDP and TCP protocols. If time permits, we will introduce these sockets in the last WH while on the distance teaching mode. This interface is commonly named after its designers, the University of California at Berkeley (Berkeley Sockets).
- **Interface to Layer 2.** Basically, these sockets allow a programmer to program against the IP protocol (At the Network layer). The technical name to these sockets is Raw Sockets and are used by notable utilities such as ping and traceroute, with which we have

---

<sup>1</sup> `send()` and `receive()` in this context must be understood as representing groups of library functions for sending and for receiving and not concrete socket library functions.

familiarity from Practice 1. The protocol family of this type of raw sockets is PF\_INET (Protocol Family\_Internet).

- **Interface to Layer 1.** At the layer 1 we find a protocol essential to this course: the Ethernet protocol. The type of socket to use is Raw Sockets again, however, in this case, the protocol family is PF\_PACKET (Protocol Family\_Packet; this protocol family is exclusive of Linux and the name, “Packet” might be considered somewhat misleading). This is the type of socket that is *the focus of this practice*. You find the location of PF\_PACKET Raw Socket interface in the Internet Architecture in Fig. 1.



**Figure 1.** Sender application on host 1 uses the PF\_PACKET socket interface to the Ethernet datalink protocol to send information to its peer application on host 2.

**Programming against PF\_PACKET socket interface.** In this practice we develop a program that sends a frame directly over the Ethernet. The program transmits the frame by calling the send() function of a PF\_PACKET/Datagram socket. The payload encapsulated into the frame results from repeatedly copying a string of bytes provided by the user on the command line. The transmission is checked by having tcpdump receive the frame, either in the same host that transmits the frame or at any other host within the same LAN.

The source code includes comments to the most important data structures and program operations. Skimming the comments should result sufficient for you to understand them; if necessary, scan the code for a more detailed study thereof. Below, an overall explanation of the program is developed focusing on the few functions that comprise it and on the socket calls and sockets data structures involved. In the upcoming WH-Practices we'll do socket programming at other layers' service interfaces and will test other sockets' features and will create slightly richer programs.

The purpose of the proposed exercises is twofold. On the one hand, they are aimed at improving the basis of your knowledge about network architecture and sockets; on the other hand, their aim is to improve your hands on ability in C programming and in Linux networking.

- The program starts with interpreting the command line arguments in function **main()**: the network interface to use for transmission, the base character string that is to be repeatedly copied into the payload and the optional total payload size, which default value is 128B.
- Function **start()** contains the core flow of operations in the program which follows this outline
  - i) Print an informational program legend
  - ii) Create a socket address structure for the raw socket that is to be created afterwards.
  - iii) Create a socket of the right type for sending directly onto an ethernet interface
  - iv) Create an array for storing the bytes that make up the payload that is to be encapsulated into the frame to transmit
  - v) Call the socket's function **sendto()** which will request the sockets layer to send the frame just built
- Sending and receiving on raw sockets entails a socket structure to store key fields from the ethernet header and other relevant information such as the network interface involved in either operation. The type of sockets used in this practice (PF\_PACKET family) employ a socket address supported by the following derived type:

```
struct sockaddr_ll
```

The postfix "ll" in the struct type name means *link level*. This struct is declared in the following header file:

```
#include <linux/if_packet.h>
```

The creation of the struct **sockaddr\_ll** necessary for sending and receiving is accomplished in function **fillSocketAddress()**. The function fills the necessary fields as in the following outline:

- The address family (PF\_PACKET)
- The interface index (ifconfig returns an interface name which must be translated into an index valid for the socket address)
- Ethernet's multiplexing key (Ethertype) to be used by this socket (Note that the selected Ethertype must be stored in *network byte order*). We use an Ethertype value that is not reserved (0x07ff).
- The remaining fields are related with ARP. We won't explain them at this time. You can see their defaults in the source code.

This function returns the struct **sockaddr\_ll** filled with the values explained above. An essential value remains to be filled at this time yet: the Destination MAC address which is set by function **setDestinationMAC()** to the broadcast address.

- The PF\_PACKET **socket** (Of type SOCK\_DGRAM) is created in line 105:
  - The socket family argument is constant PF\_PACKET
  - The socket communication style is SOCK\_DGRAM, this means that the programmer only provides the payload to be sent. The header is built by the sockets library.
  - Last argument is the Ethertype field value (The same that we created above for the socket address)
- The payload to be included in the frame to be sent over Ethernet is created in function **buildPayload()**. This function appends a number of copies of a string provided by the user on the command line. The total number of bytes to be included in the payload is given in the third command line argument and has a default value of 128 Bytes and a maximum given by the default Ethernet MTU of 1500 Bytes.
- Function **setDestinationMAC()** sets the header's Destination MAC address to the broadcast address. This program, per se, sends only to the Broadcast address.

## Sundry technical information

MAC addresses of relevant Network 2 hosts's interfaces

Host	Interface name	IP	MAC
Internal 22	eno1	192.168.86.22	e0:d5:5e:dd:ec:67
Internal 33	eno1	192.168.86.33	e0:d5:5e:d8:84:b6
Internal 44	eno1	192.168.86.44	e0:d5:5e:d8:82:26

### Source code

<http://paloalto.unileon.es/cn/Q/dgramPFPACKETSend.c>

### Source code

## Exercises for practice

1. Checking the base program:
  - a. Upload the source file to your account in paloalto with scp and compile it by issuing a command like this:
 

```
[internal 11] $ gcc -o dgramPFPACKETSend dgramPFPACKETSend.c
```
  - b. Check that the program runs; provide no command line arguments at this time. Observe the information provided by the program about the command line syntax expected.



- c. Send the program to the administrator's FIFO to have its CAP\_NET\_RAW capability set.
- d. Open a new session in paloalto.unileon.es with `ssh -p 50500`. Use that session for connecting with another host in Network 2. Power it up if necessary by sending it the magic packet (Lookup the host adapter's MAC in the table above). After waiting a couple minutes, issue an `ssh` to the chosen host and log on with administrator and password '19xxdpq16'. When you are in session, type an appropriate `tcpdump` command line that will receive the frame sent by program `dgramPF PACKETSend`; include the `tcpdump` command here alongside with the received frame:
- e. While `tcpdump` is waiting to receive the expected frames, start the `dgramPF PACKETSend` program and pass it the necessary command line arguments:
  - interface name
  - A base string like 'Andra Tutto Bene'
  - The total payload size to send

***Notice:** The following exercises entail that you modify the base program. You can modify the source code file `dgramPF PACKETSend.c` at the trampoline host in Network 2 by using the `vi` editor (Or `vim`); alternately, you can do the modifications locally in your home computer with your chosen editor and have the source code uploaded to the trampoline at every increment and then have it recompiled (Every time you modify the source code, either remotely or locally, you'll have to recompile it).*

2. Have the program accept the Destination MAC address from the command line in hexadecimal base and separating each byte with a colon (The same MAC address format used by `ifconfig`). This argument is not optional.
  - a. Create a few unit tests to check that your program functions as expected. Monitor your program's operation with `tcpdump` at another host in RemLabB6.
3. Have the program accept the Ethertype from the command line. This argument is not optional.
  - a. Create a few unit tests to check that your program functions as expected. Monitor your program's operation with `tcpdump` at another host in RemLabB6.

## Program source code

```
/*
 * Conceptual Computer Networks textbook
 * CN course 2020
 * Express Practice: Simple Datagram PF_PACKET DGRAM send program
 * dgramRawSend.c
 * All rights reserved:
 * (C) 2020 by José María Foces Morán & José María Foces Vivancos
 */

#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

#include <fcntl.h>
#include <memory.h>

#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <unistd.h>

#include <linux/if_ether.h>
#include <linux/if_packet.h>
#include <net/ethernet.h>
#include <net/if.h>

#include <signal.h>
#include <errno.h>
#include <sys/time.h>

#define byte u_char
#define TRUE 1
#define ETHERTYPE_EXPERIMENTAL 0x07ff
#define DEFAULT_MTU 1500

void setDestinationMAC(byte *p) {

    /*
     * Fill the byte array pointed to by variable p with the
     * bytes that comprise the broadcast address (All 1's)
     */
    p[0] = (byte) 0xff;
    p[1] = (byte) 0xff;
    p[2] = (byte) 0xff;
    p[3] = (byte) 0xff;
    p[4] = (byte) 0xff;
    p[5] = (byte) 0xff;

}

/*
 * Create a "large" payload of size payloadSize by cloning baseDataToClone[] as
 * many times as it fits including the remainder size
 */
char * buildPayload(char *baseDataToClone, unsigned baseDataSize, unsigned
payloadSize) {

    char *p = (char *) '\0';
```



```

    if (payloadSize == 0 || baseDataSize == 0) {
        fprintf(stderr, "Error: Requested payload size and base data size must be both
greater than 0.\n");
        return p;
    }

    if (payloadSize > DEFAULT_MTU) {
        fprintf(stderr, "Error: Requested payload size exceeds Ethernet maximum
MTU.\n");
        return p;
    }

    if (payloadSize < baseDataSize) {
        fprintf(stderr, "Error: Payload size cannot fit the base data to be
cloned.\n");
        return p;
    }

    /*
    * Request dynamic memory space for the payload to be built
    * p is a sentinel marking the start of the payload being built
    * We'll return p to our calling function
    */
    p = malloc(payloadSize);

    /*
    * Copy p into q and use the latter for indexing the payload as
    * we fill it in the for loop below
    */
    char *q = p;

    /*
    * Copy baseDataToClone as many times as it fits payloadSize
    * Each copy of baseDataSize bytes is made by calling memcpy()
    */
    for (int i = 0; i < (payloadSize / baseDataSize); i++) {
        memcpy(q, baseDataToClone, baseDataSize);
        //Move q pointer forward baseDataSize bytes
        q += baseDataSize;
    }

    /* If the integer division (payloadSize / baseDataSize) produces a
    * remainder (payloadSize % baseDataSize), copy the number of bytes
    * represented by the remainder from baseDataToClone to q
    */
    memcpy(q, baseDataToClone, payloadSize % baseDataSize);

    return p;
}

void printProgramLegend(char *payload) {

    printf("Send a frame with PF_PACKET/SOCK_DGRAM\n");
    printf("\tDMAC = ff:ff:ff:ff:ff:ff\n");
    printf("\tSMAC = Network Interface's MAC\n");
    printf("\tEthertype = %hx\n", ETHERTYPE_EXPERIMENTAL);
    //printf("\tPayload=\"%s\"", payload);

    fflush(stdout);
}

```

```

/*
 * This function fills the fields of the socket address structure
 * Some of the come from the command line arguments passed by the user
 *
 * u_int16_t is used for representing the ethertype
 * u_int16_t is declared int /usr/include/x86_64-linux-gnu/sys/types.h
 * with #include <sys/types.h>
 *
 * __be16 is defined in /usr/include/linux/types.h and is also used for
 * representing ethernet's ethertype field
 */
struct sockaddr_ll fillSocketAddress(char *ifName, u_int16_t ethertype) {

    /*
     * sockaddr_ll has slightly different used when sending than are used when
     * receiving.
     * When sending, sockaddr_ll stores the Destination MAC and the
     * multiplexing key (Ethertype) and the index of the interface to be used for
     * actually transmitting the frame
     *
     * When receiving, sockaddr_ll stores the Source MAC address, the received
     * Ethertype and the interface index the frame was received onto
     */
    struct sockaddr_ll socketAddress;

    //T
    socketAddress.sll_family = PF_PACKET;

    /* Index of network interface */
    socketAddress.sll_ifindex = if_nametoindex(ifName);
    if (socketAddress.sll_ifindex == 0) {
        perror("Error indexing interface name");
        exit(-2);
    }

    /* Address length*/
    socketAddress.sll_halen = ETH_ALEN;

    //Ethertype translated to Network Byte Order
    socketAddress.sll_protocol = htons(ethertype);

    //arp-related
    socketAddress.sll_hatype = 0;

    //arp-related
    socketAddress.sll_pkttype = 0;

    return socketAddress;
}

void start(char *ifName, char *baseDataToClone, int baseDataSize, int payloadSize) {

    /*
     * Print the frame fields when program begins to run
     */
    printProgramLegend(baseDataToClone);

    /*
     * This struct stores basic Raw socket parameters such as:
     * multiplexing key (Ethertype), Destination MAC, etc
     */

```

```

    struct sockaddr_ll socketAddress = fillSocketAddress(ifName, (u_short)
ETHERTYPE_EXPERIMENTAL);
    setDestinationMAC(&(socketAddress.sll_addr[0]));

/*
 * Create a socket with the following three actual parameter values:
 *
 * Arg 1: PF_PACKET is the address family used by Ethernet/Datalink sockets in
Linux
 *
 * Arg 2: SOCK_DGRAM is the type of communication style to be used with this
socket:
 * - SOCK_DGRAM means that the programmer is letting the building of the
 *   frame's header to the sockets layer (The service interface itself),
 *   i.e., the programmer is not building the header but is only the payload
 *   as we did above.
 *
 * - The other option available for this argument is constant SOCK_RAW
 *   which means that the programmer is providing a full datalink header
 *
 * Arg 3: Ethernet's multiplexing key (Ethertype); in this case we are using
 * - ETHERTYPE_EXPERIMENTAL (0x07ff) which is not reserved. Since we already
 *   loaded value ETHERTYPE_EXPERIMENTAL into the socket address created above,
 *   we use it again, for consistency (socketAddress.sll_protocol)
 */
int sock = socket(PF_PACKET, SOCK_DGRAM, socketAddress.sll_protocol);

/*
 * Call function to have the payload built from a base array of bytes that is
 * going to be cloned a number of times. baseDataToClone is entered by the user
 * on the command line. payloadSize is the total size of the payload
 */
byte *p = buildPayload(baseDataToClone, baseDataSize, payloadSize);
if (p == (byte *) '\0') {
    exit(-1);
}

/*
 * Finally, the data is ready to be sent onto socket sock
 * sock:      The socket created above
 * p:         Pointer to an array of bytes (unsigned char) that contains the
 *            data to be sent onto the socket
 * payloadSize:
 *            The size of the array pointed to by p in bytes
 * 0:         Options for this socket
 * socketAddress:
 *            Storage for the socket address which contains fields such as:
 *            • Dest MAC address
 *            • Ethertype (Ethernet's multiplexing key)
 *            • Consult /usr/include/linux
 */
if (sendto(sock, p, payloadSize, 0, (struct sockaddr *) &socketAddress, sizeof
(socketAddress)) == -1) {
    printf("\nsendto() call failed\n");
    perror("sendto: ");
    exit(-1);
}
}

int main(int argc, char** argv) {

/*

```

```
* Command-line processing
*/
if (argc == 3) {
    //call start with default size of 128B
    start(argv[1], argv[2], strlen(argv[2]), 128);
    printf("Simple frame successfully sent to the broadcast address via %s.\n",
argv[1]);
} else if (argc == 4) {
    int paySize = atoi(argv[3]);
    if (paySize < 0 || paySize > DEFAULT_MTU) {
        fprintf(stderr, "Invalid value for Payload size to send; should be [1,
1500]\n", argv[0]);
        exit(-1);
    }
    //call start by passing no default parameter
    start(argv[1], argv[2], strlen(argv[2]), paySize);
    printf("Simple frame successfully sent to the broadcast address via %s.\n",
argv[1]);
} else {
    fprintf(stderr, "Usage: %s\n", argv[0]);
    fprintf(stderr, "\t<Network Interface>\n");
    fprintf(stderr, "\t<Data to be cloned between quotation marks (\"...\")>\n");
    fprintf(stderr, "\t[Payload size to send (Max. 1500B, Default 128B)]\n");

    exit(-1);
}
}
```